

By Esther Ng Xin Yue

# **FEATURE ATTRIBUTION FOR MACHINE LEARNING MODELS**

**EXACT SHAP COMPUTATION ON TRACTABLE BOOLEAN CIRCUITS**

03 Feb 2026

# WHAT IS SHAP?

SHAP (SHapley Additive exPlanations) is a feature attribution method for explaining individual model predictions.

## Cooperative game theory

- **Players** → input features  $x$
- **Game payout** → model prediction  $f(x)$
- **Goal** → fair distribution of the payout amongst the players

$$f(x) = \Phi_0 + \sum_{i=1}^d \Phi_i$$

## Why is SHAP so popular?

- Model-agnostic
- Strong theoretical guarantees
- widely used in model debugging, fairness audits, and compliance testing

# WHY IS SHAP COMPUTATIONALLY HARD?

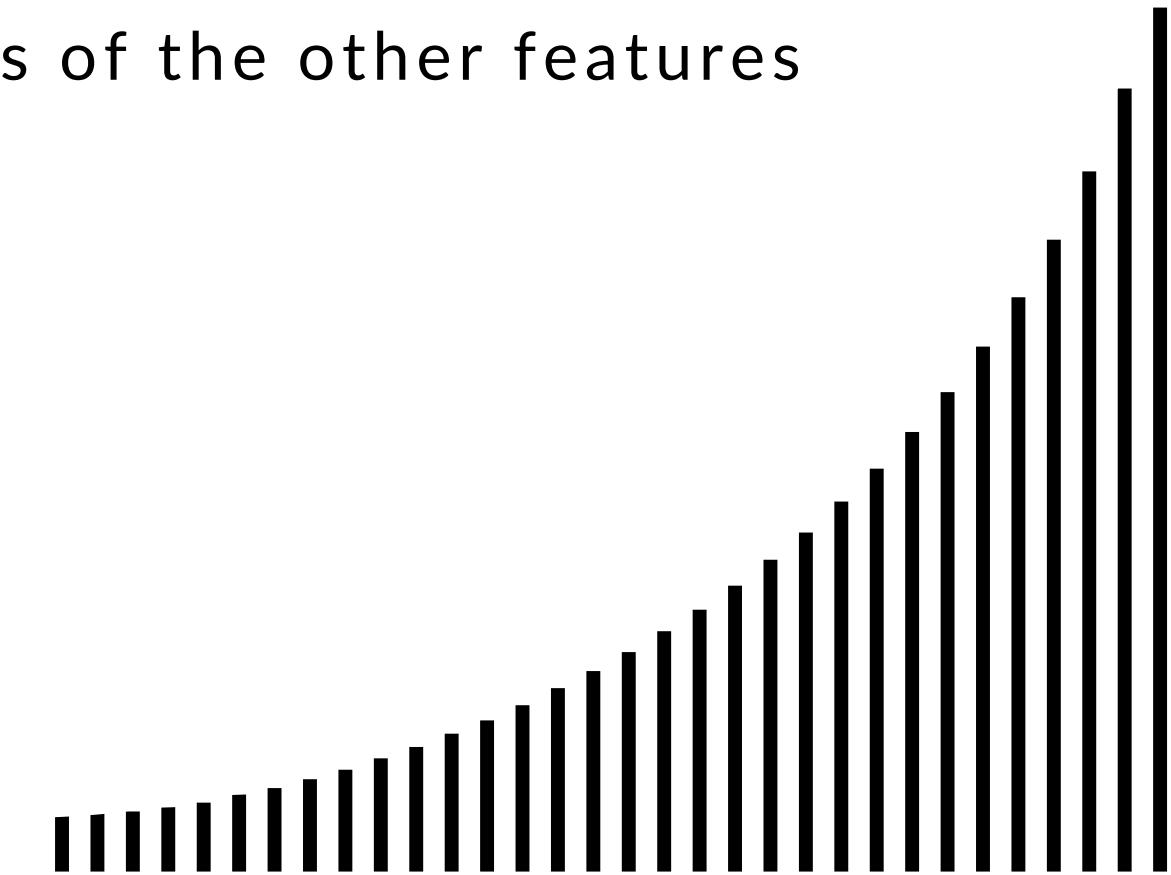
01

## EXPONENTIAL NUMBER OF FEATURE COALITIONS

- SHAP is based on Shapley values
- For each feature, SHAP considers all subsets of the other features

Number of coalitions per feature:

$$2^{n-1}$$



# WHY IS SHAP COMPUTATIONALLY HARD?

## 02 EACH COALITION REQUIRES A MODEL EVALUATION

For each subset  $S$ , SHAP evaluates a value function  $v(S)$ :

- This fixes features in  $S$
- Integrates out all remaining features
- Computes an expected model output

$$v(S) = E [f(X) | X_s = x_s]$$

# WHY IS SHAP COMPUTATIONALLY HARD?

## 03 #P-HARDNESS OF EXACT SHAP

- Computing  $v(S)$  reduces to weighted model counting
- Weighted model counting is #P-hard
  - Exact SHAP is intractable in general

# **CAN WE RESTRICT THE MODEL CLASS SO THAT SHAP BECOMES TRACTABLE?**

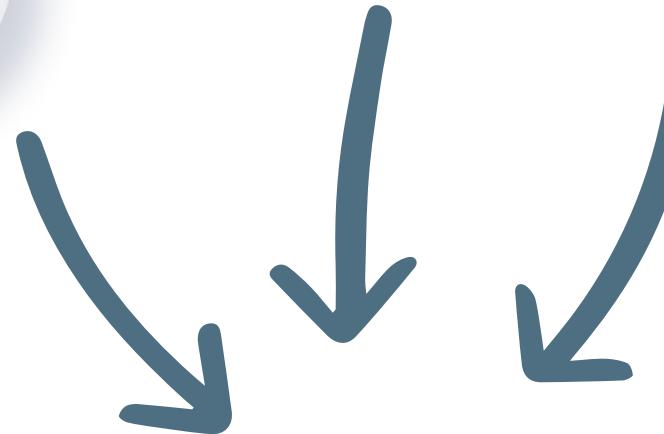
A problem is tractable if it can be solved efficiently,  
even as the problem size grows.

# YES!

with a tractable circuit representation

D-DNNFS

SENTENTIAL  
DECISION DIAGRAMS  
(SDDS)



## STRUCTURAL CONSTRAINTS

### Decomposability

- Subcircuits depend on disjoint sets of variables
- Independent parts can be evaluated separately

### Determinism

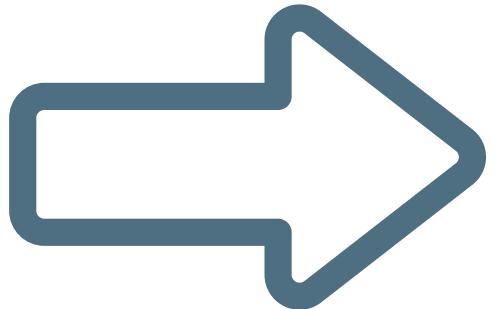
- At most one branch is true for any input
- Prevents double-counting of probability mass or explanations

### Structured Decomposability

- All decisions follow a fixed variable structure
- Ensures global consistency across the circuit

## These properties allow:

- Efficient model counting
- Efficient conditional expectations
- Efficient SHAP aggregation



## Key Insight for SHAP

- SHAP requires many conditional expectations
- In general models: #P-hard
- In tractable circuits:
  - Each conditional expectation reduces to a single linear-time circuit traversal

# WHAT ARE SDDS?

An SDD expresses a Boolean function as a disjunction of mutually exclusive cases:

$$(p_1 \cap s_1) \cup (p_2 \cap s_2) \cup \dots$$

## Key Structural Properties

### Decomposability

- primes and subs depend on disjoint variable sets

### Determinism

- no two cases can be true at the same time

### Structured Decomposability

- all decisions follow a fixed variable tree (vtree)

- Instead of checking all combinations of features,
- an SDD organises them so we never double-count or overlap work.
- supports:
  - Linear-time model counting
  - Efficient conditioning
  - Exact probability computation
  - Exact SHAP computation

SDDs avoid redundant computation by structuring decisions

# Computing SHAP Scores on SDDs

Structure → Tractable Inference → Tractable SHAP

# SHAP Component

# SDD Interpretation

FEATURE SUBSET S



PARTIAL ASSIGNMENT OF  
VARIABLES

$$X_s = x_s$$

VALUE FUNCTION OF V(S)



EXPECTED OUTPUT UNDER  
CONDITIONING ON S

$$E [f(X) | X_s = x_s]$$

MARGINAL CONTRIBUTION



DIFFERENCE OF TWO  
CONDITIONAL EXPECTATIONS

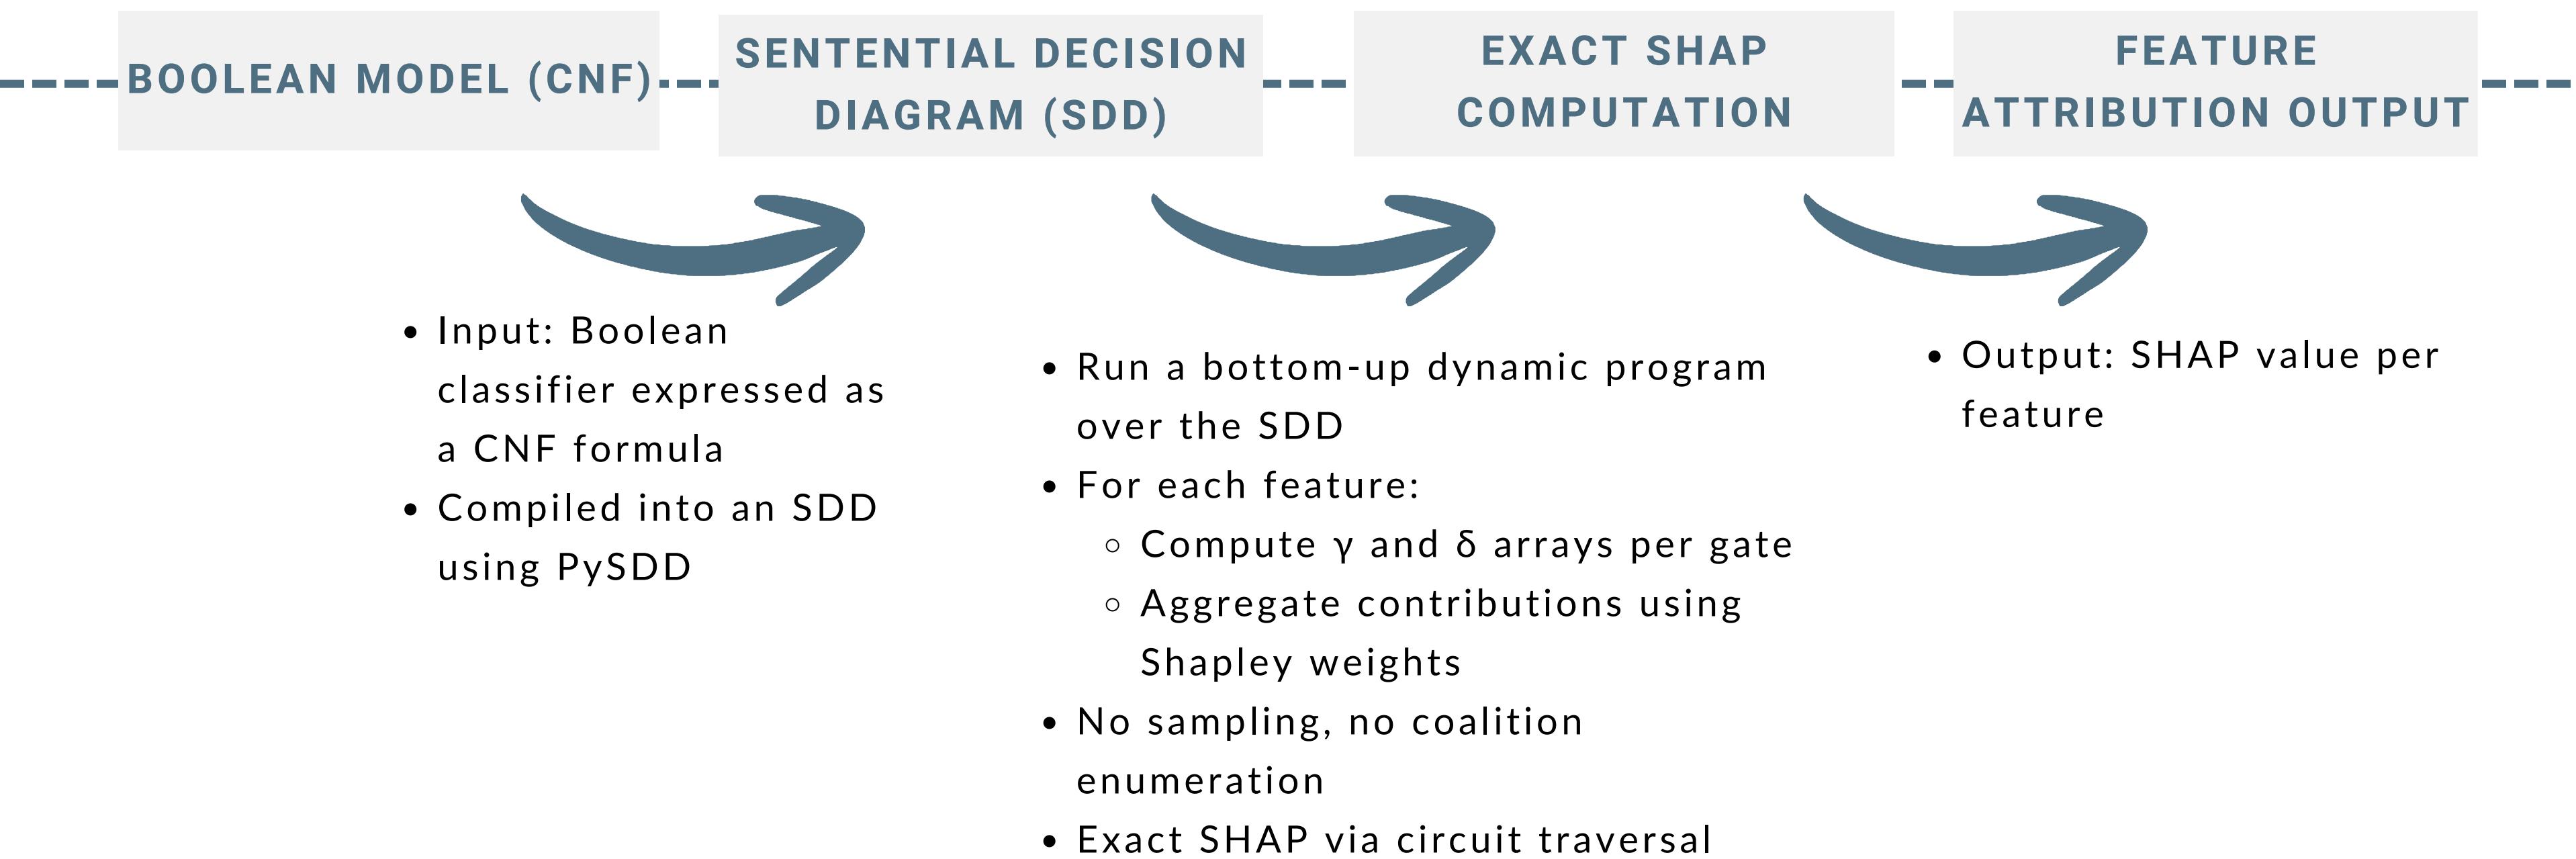
$$E [f(X) | x_s = 1] - E [f(X) | x_s = 0]$$

SUMMATION OVER COALITION



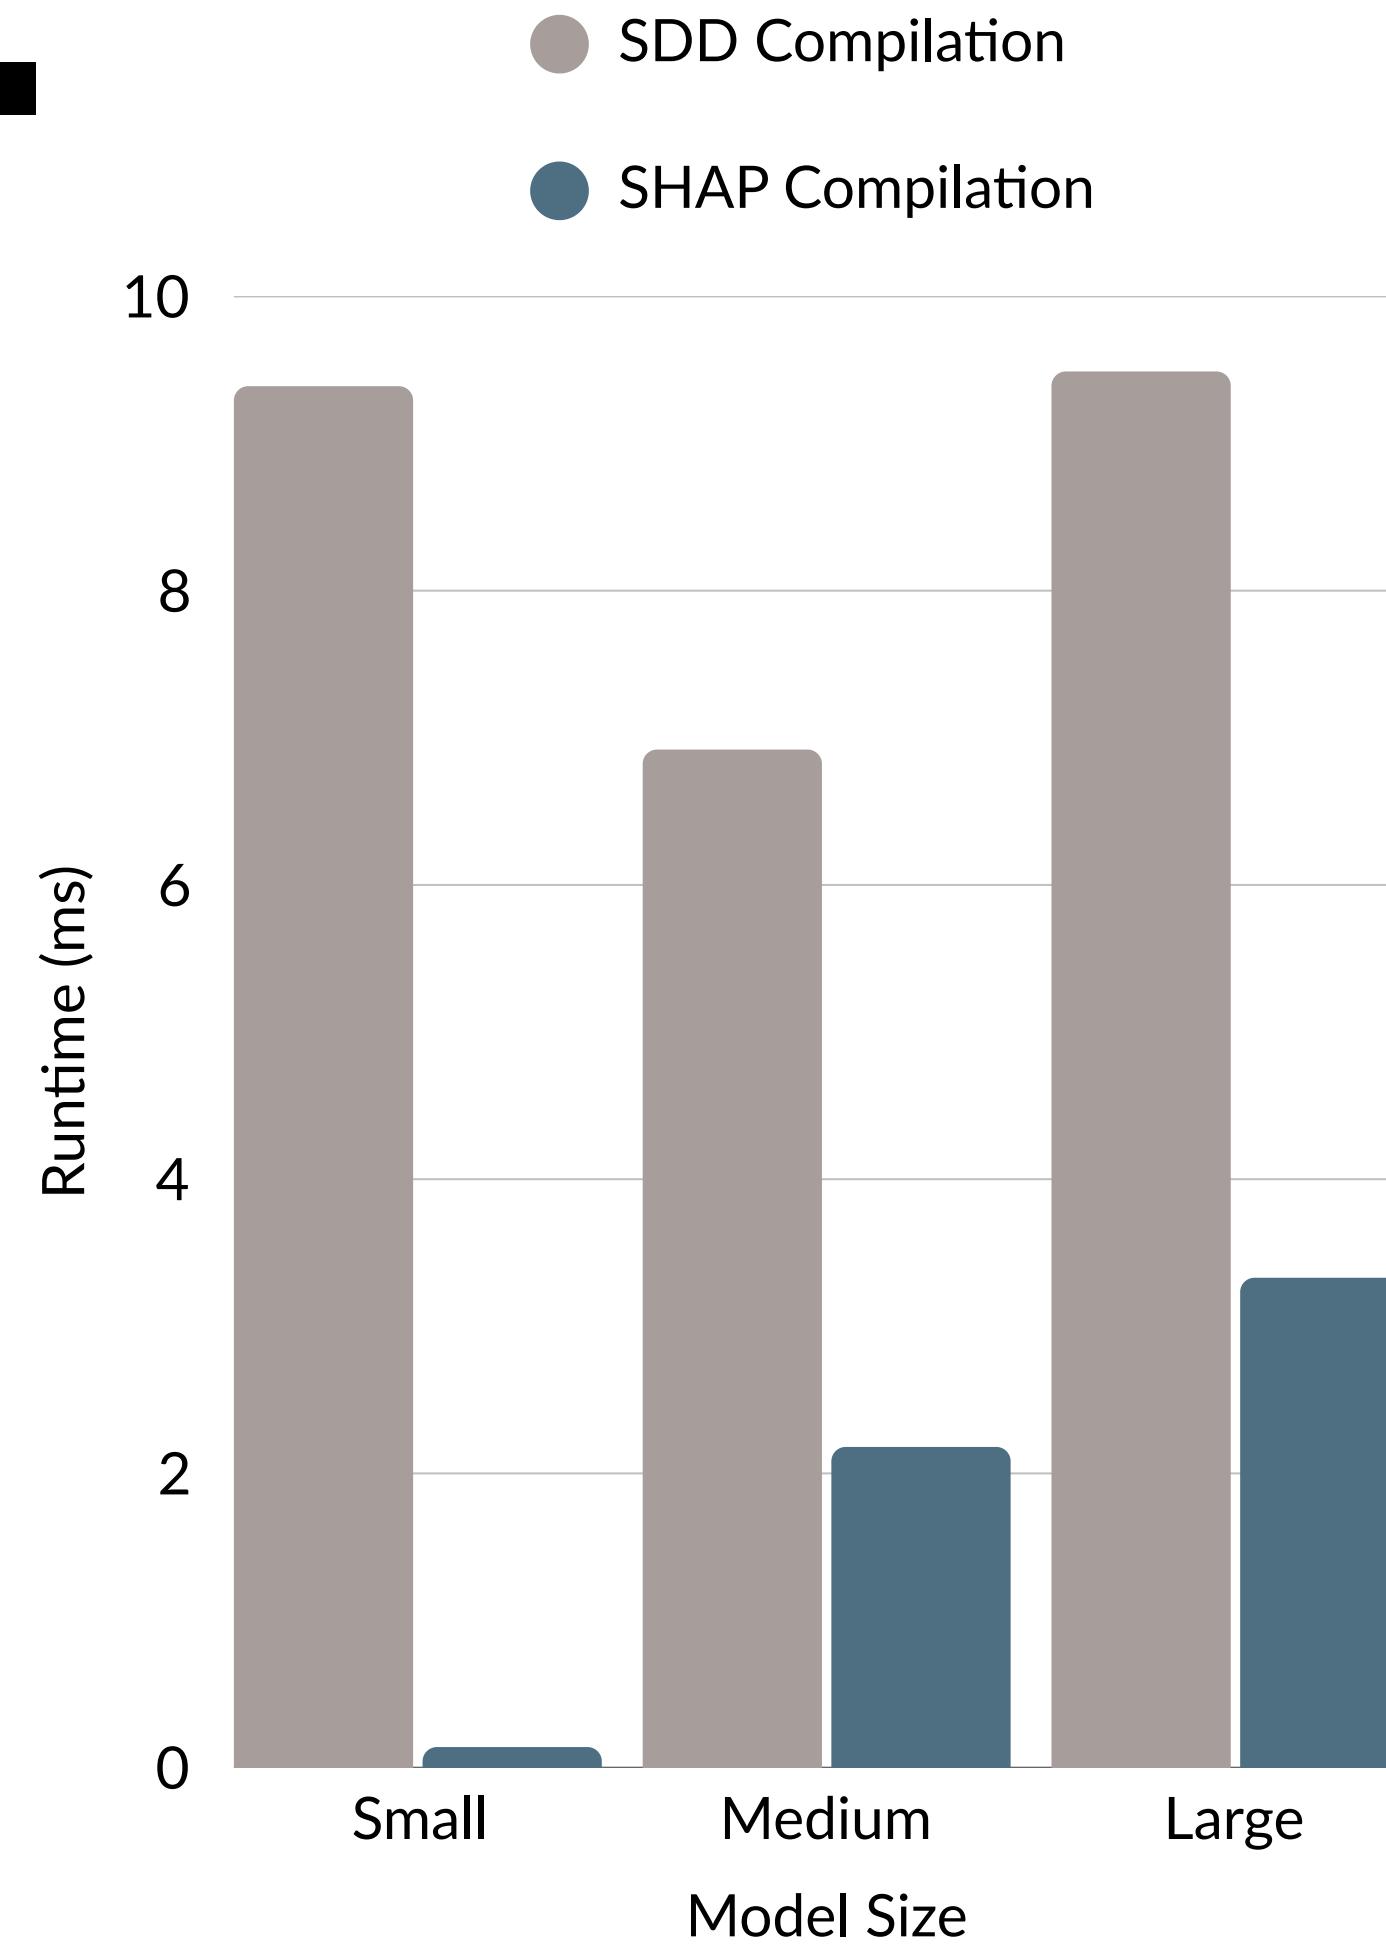
AGGREGATION VIA DYNAMIC  
PROGRAMMING

# SYSTEM ARCHITECTURE



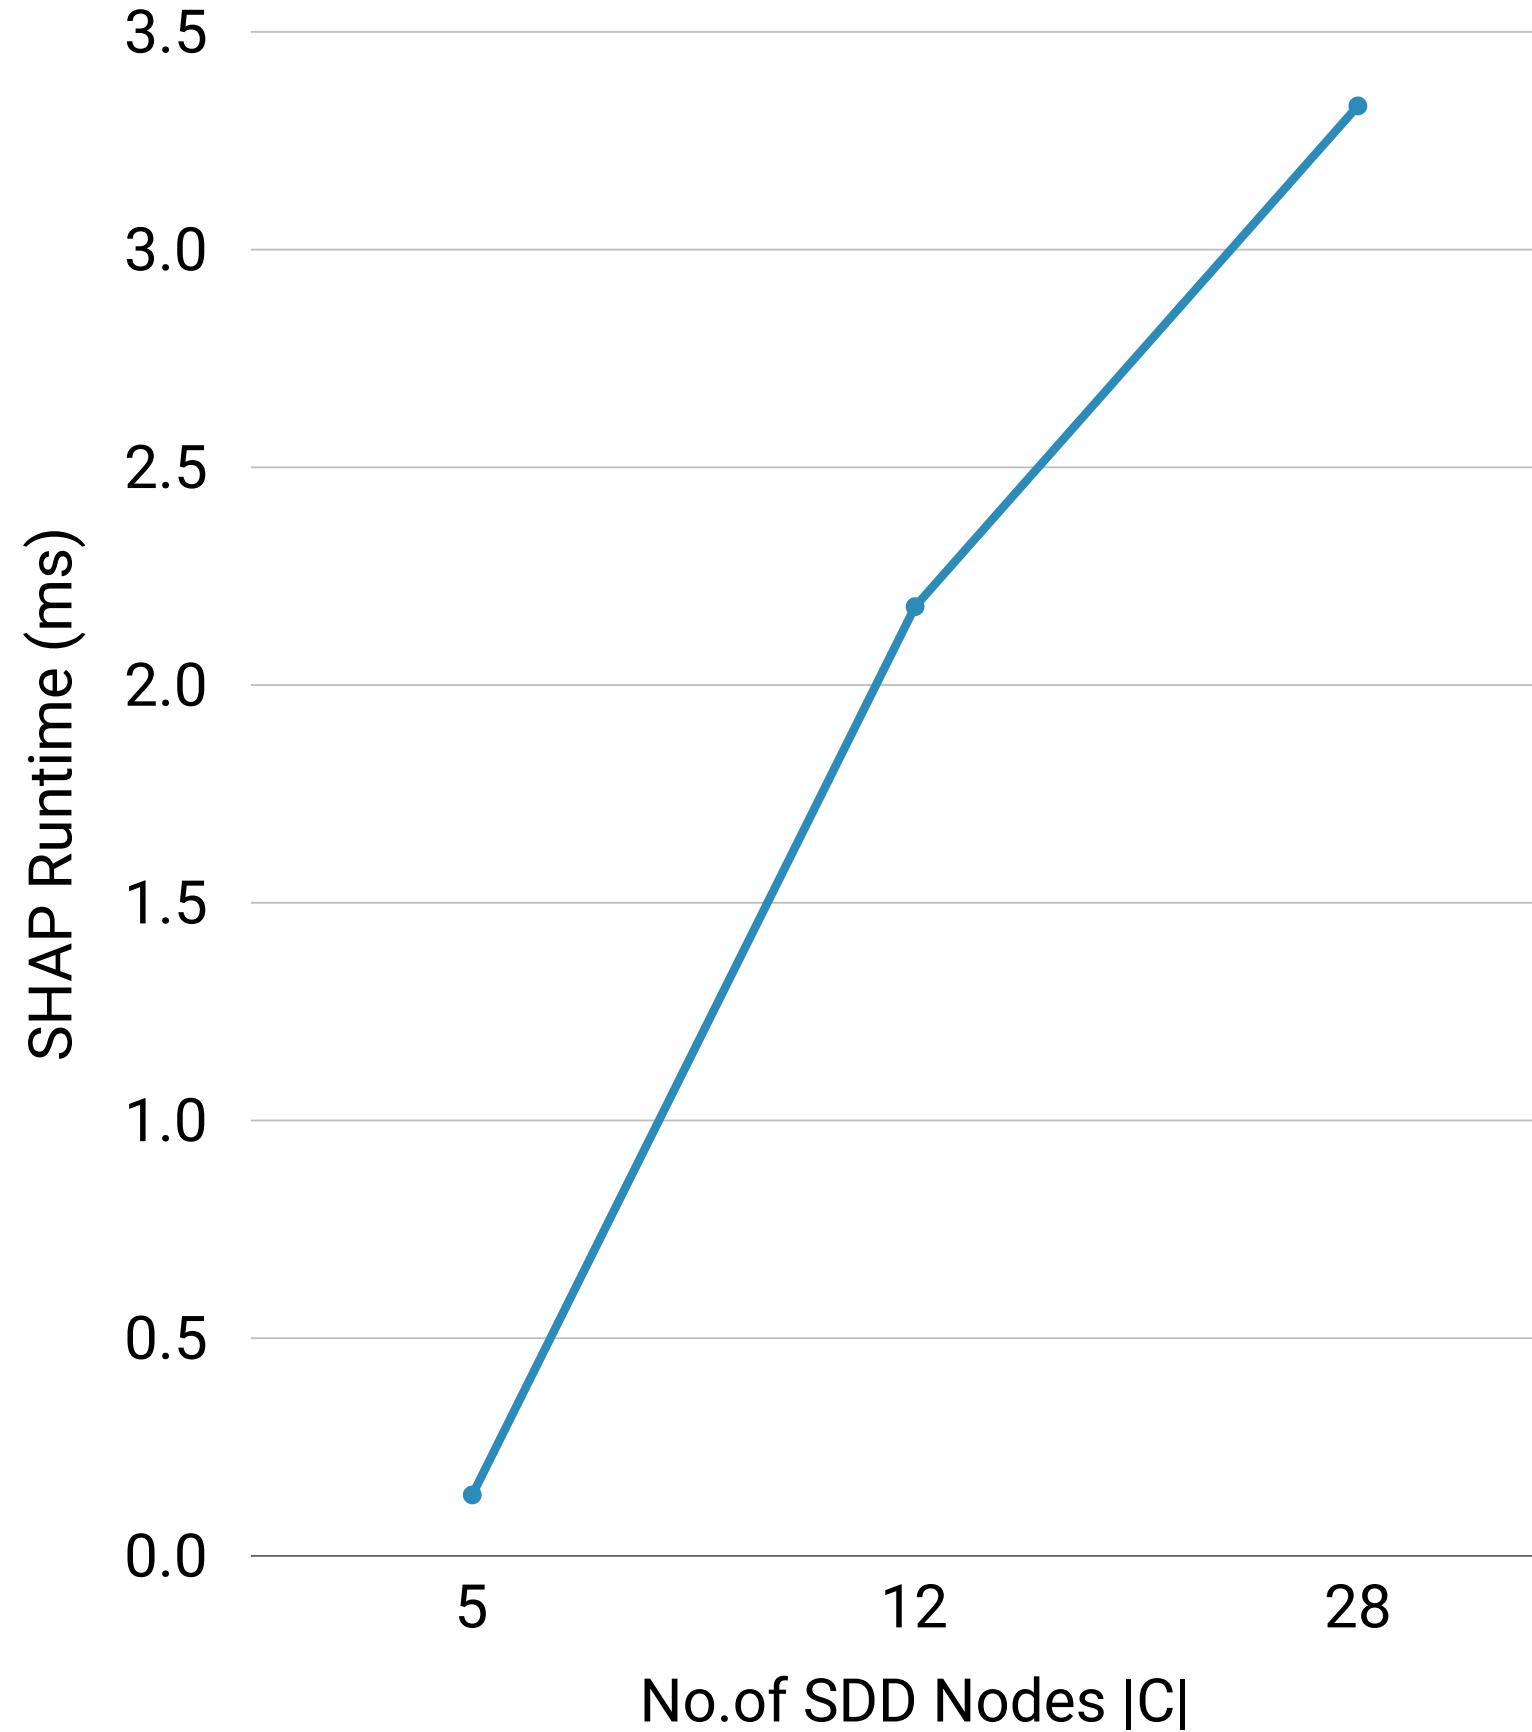
# BENCHMARKING

BENCHMARK 1: Compilation vs SHAP runtime (means, ms)			
Test Case	Compile(ms)	SHAP(ms)	Total(ms)
Small (2 vars, 4 clauses) - Symmetric	9.39	0.14	9.53
Medium (4 vars, 2 clauses) - Independent OR	6.92	2.18	9.11
Large (5 vars, 10 clauses) - Constraint SAT	9.49	3.33	12.86
BENCHMARK 2: Runtime vs circuit size (SDD nodes) [mean_shap]			
Test Case	SDD(nodes)	Runtime(ms)	
Small (2 vars, 4 clauses) - Symmetric	5	0.14	
Medium (4 vars, 2 clauses) - Independent OR	12	2.18	
Large (5 vars, 10 clauses) - Constraint SAT	28	3.33	
BENCHMARK 3: Compilation vs SHAP cost breakdown (mean % of total)			
Test Case	Compile%	SHAP%	(rest)
Small (2 vars, 4 clauses) - Symmetric	98.5	1.5	0.1
Medium (4 vars, 2 clauses) - Independent OR	76.0	23.9	0.1
Large (5 vars, 10 clauses) - Constraint SAT	73.9	25.9	0.2
BENCHMARK 4: SHAP runtime vs number of variables (means)			
Test Case	Vars	SHAP(ms)	
Small (2 vars, 4 clauses) - Symmetric	2	0.14	
Medium (4 vars, 2 clauses) - Independent OR	4	2.18	
Large (5 vars, 10 clauses) - Constraint SAT	6	3.33	



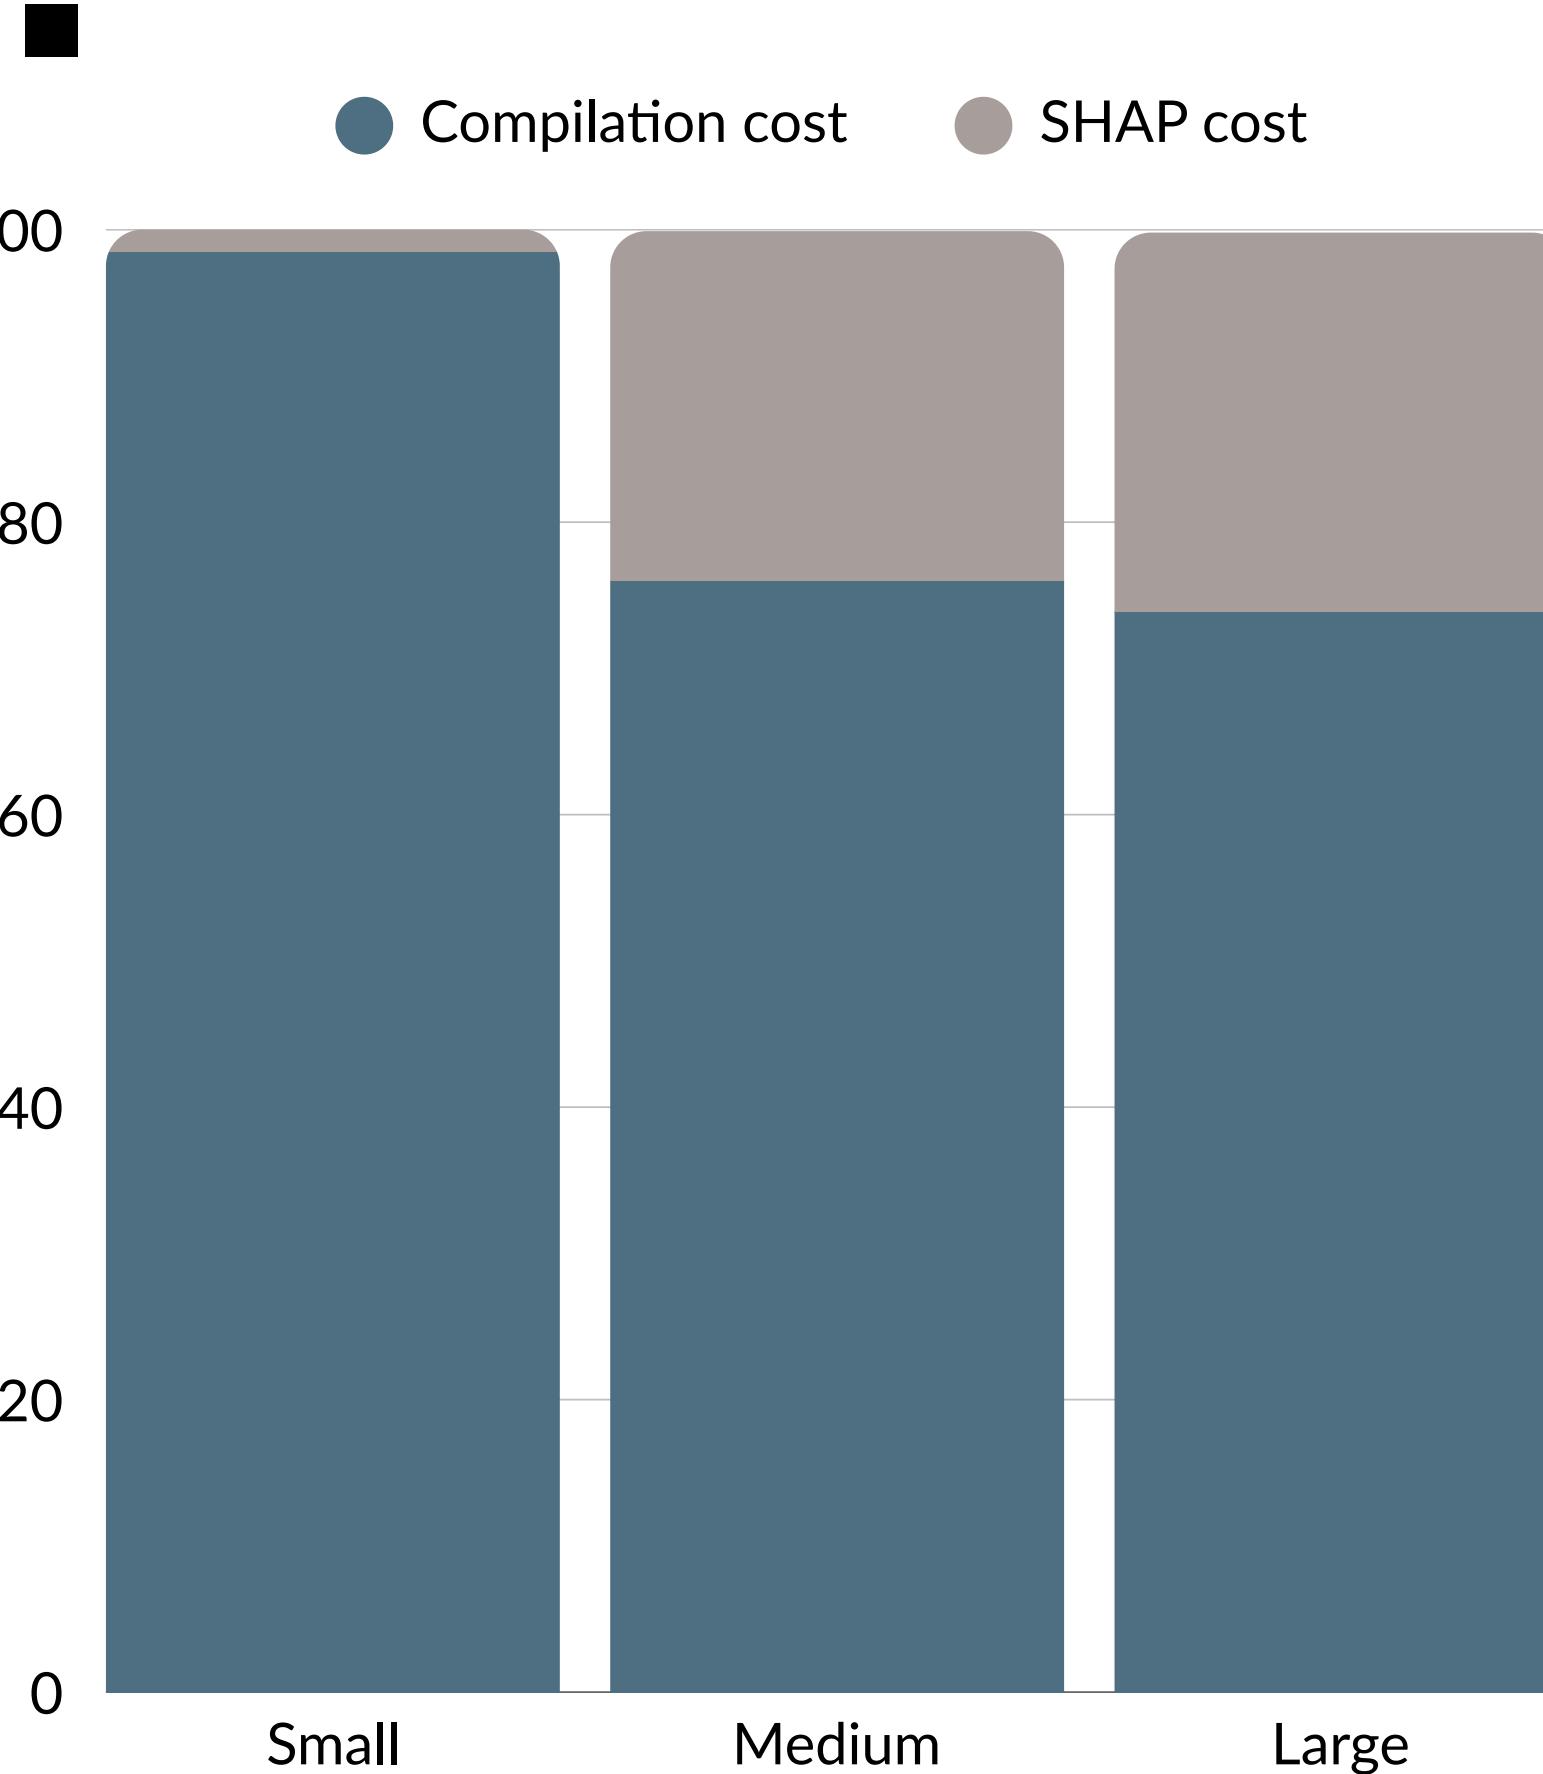
## Compilation vs SHAP Runtime

- Compilation dominates total runtime
- SHAP computation is consistently small
- Once compiled, explanations are cheap



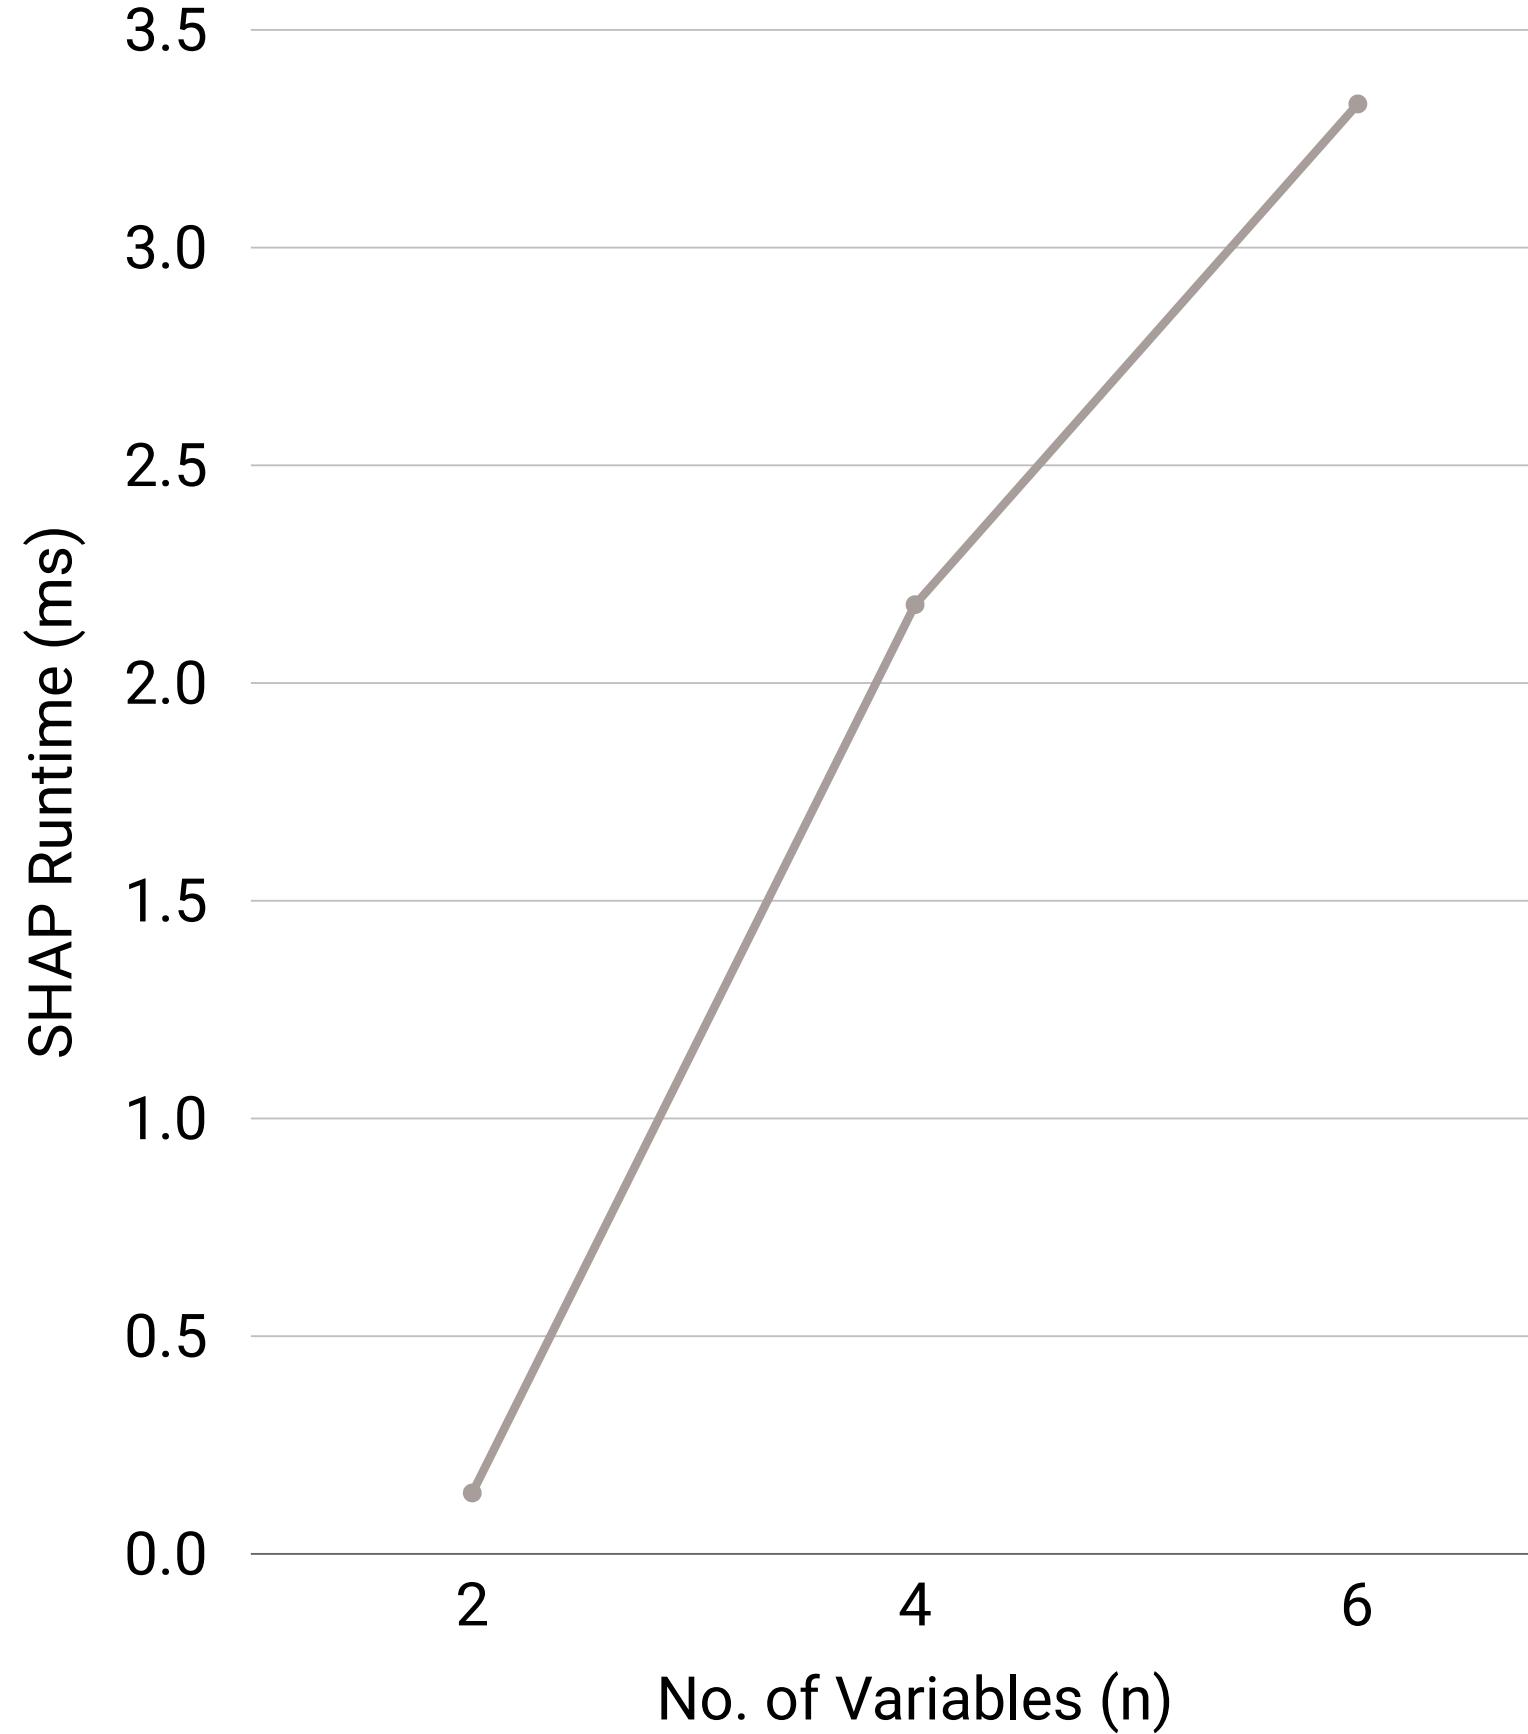
## Runtime vs Circuit Size

- Smooth, near-linear scaling
- No exponential blow-up
- Matches theoretical complexity



## Compilation vs SHAP Cost Breakdown

- SHAP is a minor fraction of total cost
- Confirms tractability in practice



## SHAP Runtime vs Number of Variables

- Polynomial growth
- No combinatorial explosion
- Practical even as n increases

# APPLICATIONS OF TRACTABLE SHAP

SHAP (SHapley Additive exPlanations) is a feature attribution method for explaining individual model predictions.

## Content Moderation Rules

- Exact attribution of which rule conditions triggered a decision
- Clear explanation for:
  - Appeals
  - Policy audits
  - Internal debugging
- “This post was flagged primarily due to X, not Y.”

## Policy Based Classifiers

- Transparent justification for accept/reject outcomes
- Stable explanations across runs
- No sampling noise
- Especially important in regulated environments.

## Hybrid AI Systems

- Feature-level explanations of the policy layer
- Clear separation between:
  - Learned behaviour
  - Enforced constraints

01

# Main Bottleneck : **SDD Compilation**

In practice: explainability is cheap once the model is compiled

# 02

# Model Class Restrictions

Only applies for models that can be expressed as boolean logic,  
and not directly applicable to large neural networks, or  
unstructured continuous models

03

# **Scalability depends on the Structure**

Worst-case SDD size can be exponential

04

# Assumptions

We are assuming fully factorised input distributions, more complex dependencies would require richer circuit representations, such as PSDDs.

# Future Research Directions

## From SDDs to Probabilistic SDDs (PSDDs)

Many real systems are inherently probabilistic

Extend exact SHAP computation from:  
Boolean expectations → probabilistic expectations

## SHAP for Bayesian Network Classifiers via SDDs

Can exact SHAP be computed efficiently for Bayesian classifiers through circuit compilation?